

Notes on Linear Sorting

Geoffrey Matthews

May 9, 2016

Comparison sorts

- The only operation that may be used to gain information about a sequence is comparisons between pairs of elements.
- All sorts seen so far are comparison sorts:
 - insertion sort
 - merge sort
 - quicksort
 - heapsort


Lower bounds for comparison sorts

- $\Omega(n)$ to examine all the input
- All sorts seen so far are $\Omega(n \lg n)$
- We will show that all comparison sorts must be $\Omega(n \lg n)$

Decision tree


- Abstraction of any comparison sort
- Represents comparisons made by
 - a specific sorting algorithm
 - on inputs of a given size
- Abstracts away everything else: control and data movement.
- We're counting *only* comparisons.

Insertion sort on three elements



- Internal nodes labeled by comparisons (original positions).
- Leaf nodes labeled by permutation of order from original.
- Number of leaves $\geq n!$.

For any comparison sort



- 1 tree for each n
- View the tree as if the algorithm splits in two at each node.
- The tree models all possible execution traces.

What is the longest path from root to leaf?

- Depends on the algorithm.
- Insertion sort: $\Theta(n^2)$
- Merge sort: $\Theta(n \lg n)$

Lemma: any binary tree of height h has $\leq 2^h$ leaves.

- $\ell = \#$ of leaves
- $h = \text{height}$
- then $\ell \leq 2^h$

Proof by induction on h :

Base: $h = 0$. Tree is just one node, which is a leaf. $1 \leq 2^h$.

Inductive step: Assume true for $h - 1$. Extend tree with as many new leaves as possible. Each leaf becomes the parent of two new leaves.

$$\begin{aligned}\# \text{ of leaves for } h &= 2(\# \text{ of leaves for } h - 1) \\ &\leq 2(2^{h-1}) \\ &= 2^h\end{aligned}$$

Theorem: any decision tree that sorts n elements has height $\Omega(n \lg n)$

- $\ell \geq n!$
- $n! \leq \ell \leq 2^h$
- $h \geq \lg(n!)$
- Sterling's approximation: $n! > (n/e)^n$
- Therefore:

$$\begin{aligned} h &\geq \lg(n!) \\ &\geq \lg(n/e)^n \\ &= n \lg(n/e) \\ &= n \lg n - n \lg e \\ &= \Omega(n \lg n) \end{aligned}$$

Sorting in linear time

- Impossible with any comparison sort.
- **Counting sort**
 - Key assumption: numbers to be sorted are integers in $\{0, \dots, k\}$.

Input: $A[1..n]$ where $A[j] \in \{0, \dots, k\}$

Output: $B[1..n]$, sorted.

Auxiliary storage: $C[0..k]$

Counting sort example

COUNTING-SORT(A, B, n, k)

```
1  let  $C[0..k]$  be a new array
2  for  $i = 0$  to  $k$ 
3       $C[i] = 0$ 
4  for  $j = 1$  to  $n$ 
5       $C[A[j]] = C[A[j]] + 1$ 
6  for  $i = 1$  to  $k$ 
7       $C[i] = C[i] + C[i - 1]$ 
8  for  $j = n$  downto 1
9       $B[C[A[j]]] = A[j]$ 
10      $C[A[j]] = C[A[j]] - 1$ 
```

A:

2 ₁	5 ₁	3 ₁	0 ₁	2 ₂	3 ₂	0 ₂	3 ₃
----------------	----------------	----------------	----------------	----------------	----------------	----------------	----------------

After second **for** loop:

C:

2	0	2	3	0	1
---	---	---	---	---	---

After third **for** loop:

C:

2	2	4	7	7	8
---	---	---	---	---	---

Counting sort example

COUNTING-SORT(A, B, n, k)

```

1  let  $C[0..k]$  be a new array
2  for  $i = 0$  to  $k$ 
3       $C[i] = 0$ 
4  for  $j = 1$  to  $n$ 
5       $C[A[j]] = C[A[j]] + 1$ 
6  for  $i = 1$  to  $k$ 
7       $C[i] = C[i] + C[i - 1]$ 
8  for  $j = n$  downto 1
9       $B[C[A[j]]] = A[j]$ 
10      $C[A[j]] = C[A[j]] - 1$ 

```

A:

2_1	5_1	3_1	0_1	2_2	3_2	0_2	3_3
-------	-------	-------	-------	-------	-------	-------	-------

C:

2	2	4	7	7	8
---	---	---	---	---	---

B:

						3_3	
	0_2					3_3	
	0_2				3_2	3_3	
	0_2		2_2		3_2	3_3	
0_1	0_2		2_2		3_2	3_3	
0_1	0_2		2_2	3_1	3_2	3_3	
0_1	0_2		2_2	3_1	3_2	3_3	5_1
0_1	0_2	2_1	2_2	3_1	3_2	3_3	5_1

Counting sort is **stable**:

- Keys with the same value appear in the same order in output as in input.

Counting sort analysis


COUNTING-SORT(A, B, n, k)

```
1  let  $C[0..k]$  be a new array
2  for  $i = 0$  to  $k$ 
3       $C[i] = 0$ 
4  for  $j = 1$  to  $n$ 
5       $C[A[j]] = C[A[j]] + 1$ 
6  for  $i = 1$  to  $k$ 
7       $C[i] = C[i] + C[i - 1]$ 
8  for  $j = n$  downto 1
9       $B[C[A[j]]] = A[j]$ 
10      $C[A[j]] = C[A[j]] - 1$ 
```

- $\Theta(n + k)$
 - which is $\Theta(n)$ if $k = O(n)$.
- How big a k is practical?
 - 32-bit values? No.
 - 16-bit? Probably not.
 - 8-bit? Maybe, depending on n .
 - 4-bit? Unless n is really small.

- Counting sort will be used in radix sort.

Radix sort example



Radix sort

- IBM in early 20th century.
- Punch card sorting machines only sorted on one column.
- Humans would reload the cards and change the column.
- Human-machine cyborg algorithm!

Key idea:

Sort *least* significant digits first.



Radix sort

RADIX-SORT(A, d)

```
1  for  $i = 1$  to  $d$ 
2      use a stable sort to sort  $A$  on digit  $i$ 
```


Radix sort correctness

- Induction on number of passes.
- Assume digits $1, \dots, i - 1$ are sorted.
- Show that a stable sort on i leaves $1, \dots, i - 1$ sorted:
 - If 2 digits in position i are different,
 - * ordering by i is correct and positions $1, \dots, i - 1$ are irrelevant.
 - If 2 digits in position i are equal,
 - * numbers are already sorted by inductive hypothesis. Stable sort leaves them that way.

Radix sort analysis

Assume we use counting sort on each digit.

- $\Theta(n + k)$ per digit
- d digits
- $\Theta(d(n + k))$ total
- If $k = O(n)$, time = $\Theta(dn)$.

Radix sort: How to break each key into digits?

- n words
- b bits/word
- Break into r -bit digits. $d = \lceil b/r \rceil$
- Use counting sort, $k = 2^r - 1$.

Example: 32-bit words, 8-bit digits.

$$b = 32$$

$$d = \lceil 32/8 \rceil = 4$$

$$r = 8$$

$$k = 2^8 - 1 = 255$$

- Time = $\Theta\left(\frac{b}{r}(n + 2^r)\right)$

How to choose r ?

- Time = $\Theta\left(\frac{b}{r}(n + 2^r)\right)$
- Balance b/r and $n + 2^r$.
- Choosing $r \approx \lg n$ gives

$$\Theta\left(\frac{b}{\lg n}(n + n)\right) = \Theta(bn/\lg n)$$

- If we choose $r < \lg n$ then $b/r > b/\lg n$ and $n + 2^r$ doesn't improve.
- If we choose $r > \lg n$ then $n + 2^r$ term gets big.
- Sort 2^{16} 32-bit numbers, use $r = \lg 2^{16} = 16$ bits. $\lceil b/r \rceil = 2$ passes.

Compare radix to merge and quick

- 1 million (2^{20}) 32-bit integers.
- Radix sort: $\lceil 32/20 \rceil = 2$ passes.
- Merge/quick: $\lg n = 20$ passes.
- Each radix “pass” is 2 passes:
 - one to take census
 - one to move data

How does radix sort violate the $\Omega(n \lg n)$ speed limit?

- Counting sort allows us to gain information about keys
 - other than by directly comparing 2 keys.
- Used keys as array indices,
 - thus getting far more information out of each key.
 - branching factor of the decision tree is k