

# Notes on Red-black Trees

Geoffrey Matthews

May 25, 2016

## Red-black trees

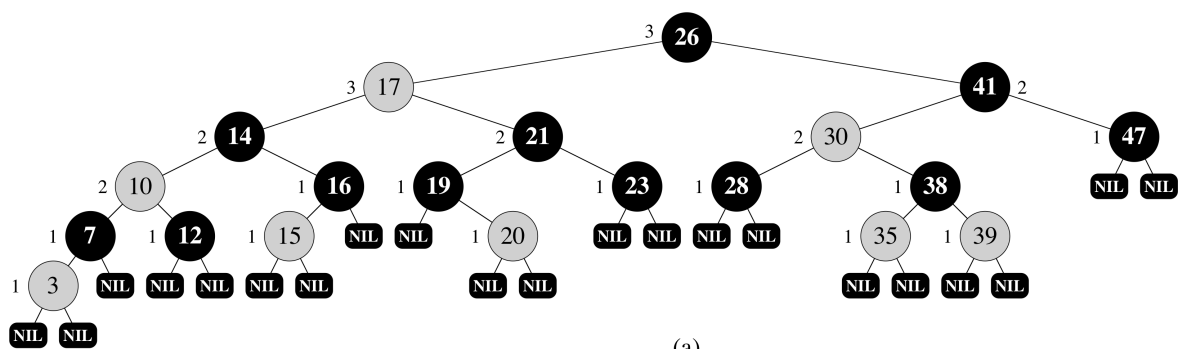
- A variation of binary search trees.
- **Balanced:** height is  $O(\lg n)$ , where  $n$  is number of nodes.
- Operations will take  $O(\lg n)$  in worst case.

## Red-black trees

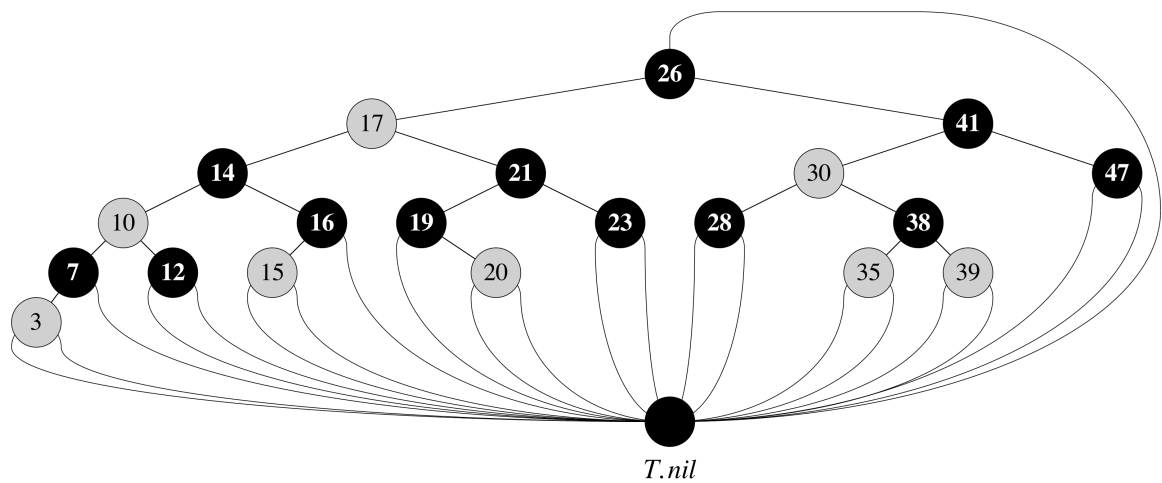
- A **red-black tree** is a binary search tree.
- One bit per node stores an attribute *color*, red or black.
- All leaves are empty (nil) and colored black.
- We use a sentinel  $T.nil$  for all the leaves of a red-black tree  $T$ .
- $T.nil.color$  is black.
- The root's parent is also  $T.nil$ .

## Red-black tree properties

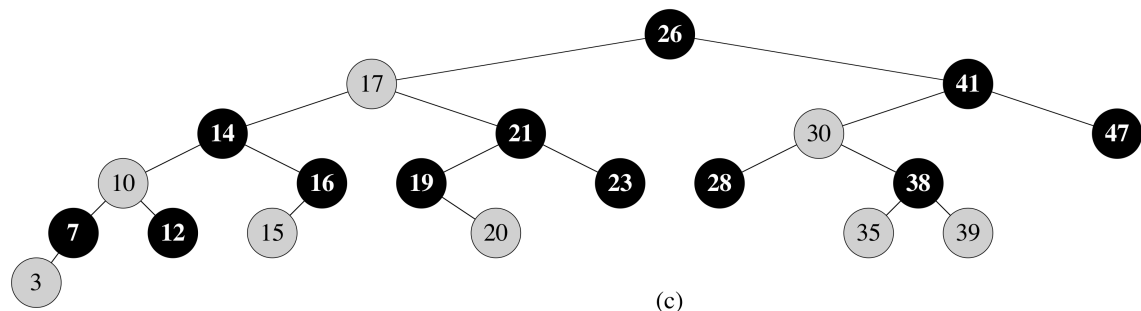
1. Every node is either red or black.
2. The root is black.
3. Every leaf ( $T.nil$ ) is black.
4. If a node is red, then both its children are black
  - Hence no two reds in a row.
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes.



(a)



(b)



(c)

## Height of a red-black tree

- **Height of a node** is the number of edges in longest path to leaf.
- **Black-height** of a node  $x$ :  $bh(x)$  is the number of black nodes (including  $T.nil$ ) on a path from  $x$  to a leaf, not counting  $x$ .
  - By property 5, black-height is well defined.
  - Changing the color of a node does not change its black-height.
  - Changing the color of a node will change the black-height of its ancestors.

**Claim 1:** Any node with height  $h$  has black-height  $\geq h/2$ .

**Proof**

- By property 4,  $\leq h/2$  nodes on the path from node to a leaf are red.
- Hence  $\geq h/2$  are black.

**Claim 2:** The subtree rooted at  $x$  contains  $\geq 2^{bh(x)} - 1$  internal nodes.

**Proof.** By induction on height of  $x$ .

**Basis:** Height of  $x = 0 \Rightarrow x$  is a leaf and so  $bh(x) = 0$ ,  $2^0 - 1 = 0$ .

**Inductive step:**

- Let the height of  $x$  be  $h$ .
- Any child of  $x$  has height  $h - 1$  and black-height either  $bh(x)$  (if the child is red) or  $bh(x) - 1$  (if the child is black).
- By inductive hypothesis, each child has  $\geq 2^{bh(x)-1} - 1$  internal nodes.
- Thus, the subtree rooted at  $x$  contains  $\geq 2 \cdot (2^{bh(x)-1} - 1) + 1 = 2^{bh(x)} - 1$  internal nodes.



**Lemma:** A red-black tree with  $n$  internal nodes and height  $h$  has

$$h \leq 2 \lg(n + 1)$$

- Recall proven claims:

- Any node with height  $h$  has black-height  $\geq h/2$ .
- The subtree rooted at any node  $x$  contains  $\geq 2^{bh(x)} - 1$  internal nodes.

**Proof**

Let  $h$  and  $b$  be the height and black-height of the root, respectively.

By the above two claims,

$$n \geq 2^b - 1 \geq 2^{h/2} - 1$$

Adding 1 to both sides and then taking logs gives

$$\lg(n + 1) \geq h/2$$

which implies that

$$h \leq 2 \lg(n + 1)$$

## Operations on red-black trees


- MINIMUM, MAXIMUM, SUCCESSOR, PREDECESSOR and SEARCH all run in  $O(h) = O(\lg n)$  time.
- INSERT, what color to make the new node?
  - Red? Might violate property 4.
  - Black? Might violate property 5.
- DELETE, what color was the old node?
  - Red? OK.
    - \* Unless successor is black?
  - Black? Could cause two reds in a row, and violate properties 2 and 5.

## Rotations

- Only pointers are changed.
- Won't upset binary-search-tree property.
- Doesn't care about red-black.

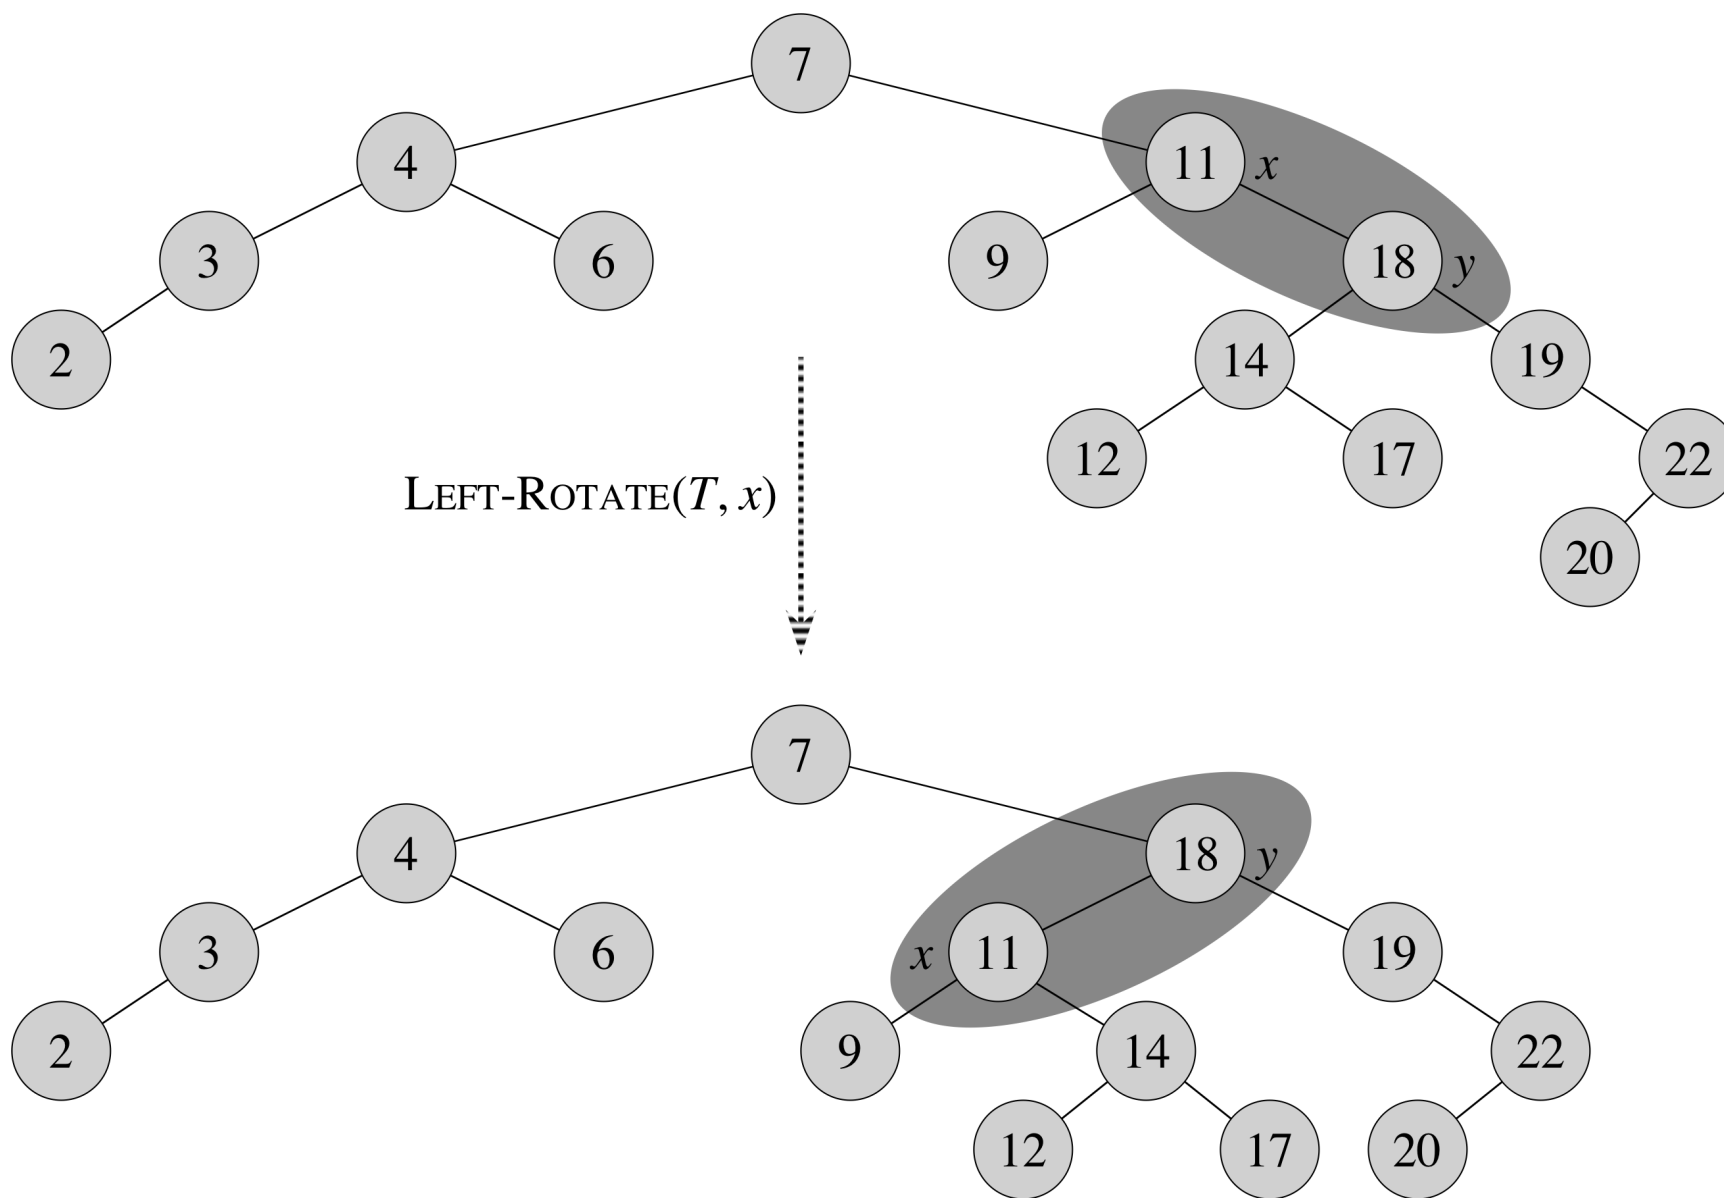
LEFT-ROTATE( $T, x$ )

```
1   $y = x.right$ 
2   $x.right = y.left$ 
3  if  $y.left \neq T.nil$ 
4       $y.left.p = x$ 
5   $y.p = x.p$ 
6  if  $x.p == T.nil$ 
7       $T.root = y$ 
8  elseif  $x == x.p.left$ 
9       $x.p.left = y$ 
10 else  $x.p.right = y$ 
11  $y.left = x$ 
12  $x.p = y$ 
```



Assumes

- $x.right \neq T.nil$
- root's parent is  $T.nil$



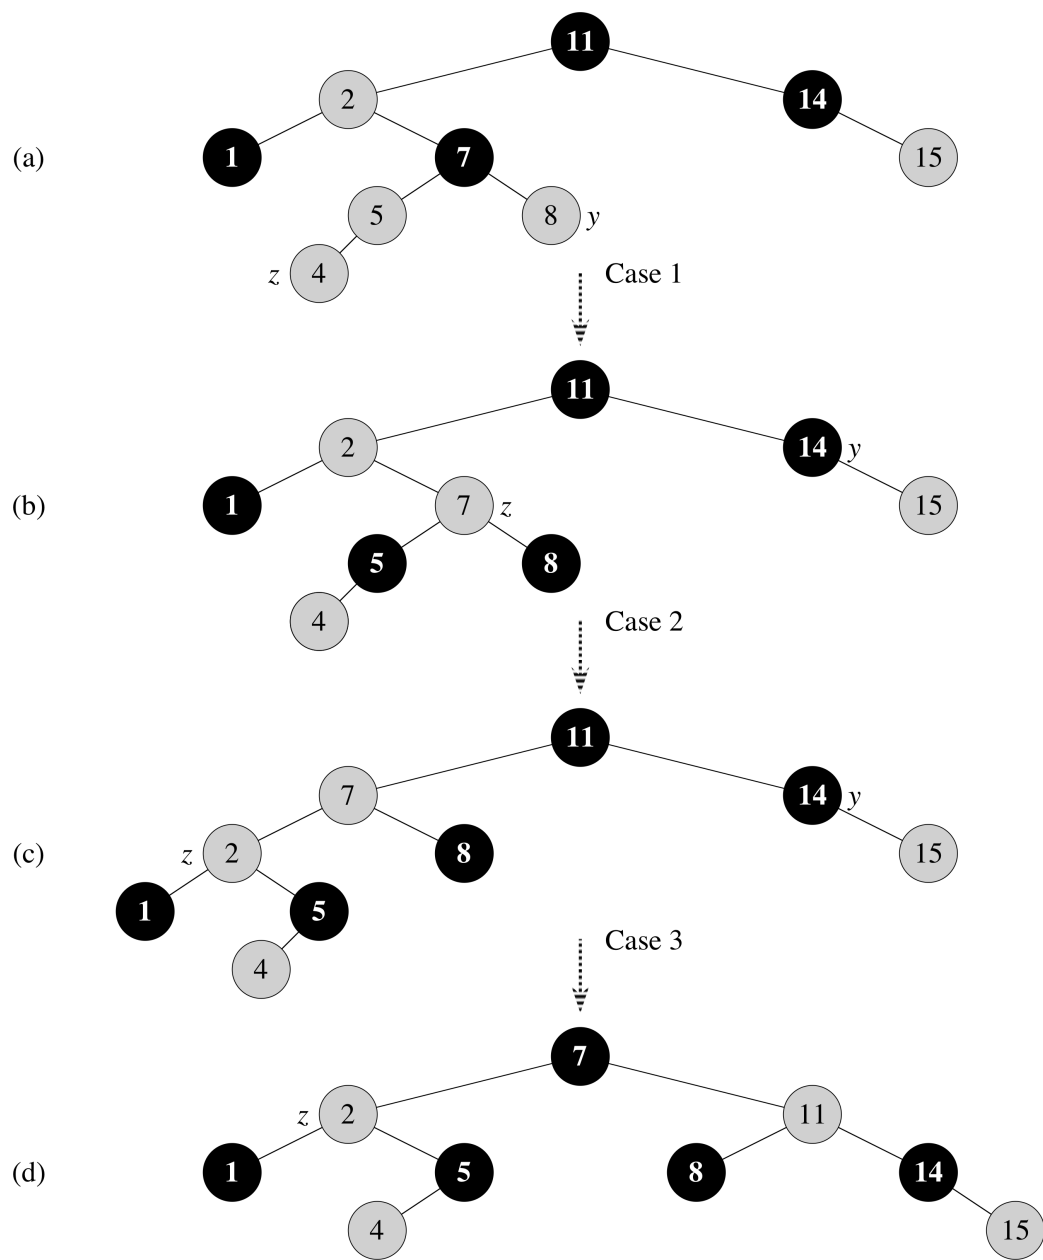
## Insertions

- Start by doing regular binary-tree insertion.
- Color new node red.
- May violate red-black tree properties:
  1. Every node is either red or black.
    - OK.
  2. The root is black.
    - New node might be root.
  3. Every leaf ( $T.nil$ ) is black.
    - OK.
  4. If a node is red, then both its children are black.
    - New node's parent might be red.
  5. For each node, all paths from the node to descendant leaves contain the same number of black nodes.
    - OK.

```

RB-INSERT-FIXUP( $T, z$ )
1  while  $z.p.color == \text{RED}$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$ 
4          if  $y.color == \text{RED}$ 
5               $z.p.color = \text{BLACK}$  // case 1
6               $y.color = \text{BLACK}$ 
7               $z.p.p.color = \text{RED}$ 
8               $z = z.p.p$ 
9          else if  $z == z.p.right$ 
10              $z = z.p$  // case 2
11             LEFT-ROTATE( $T, z$ )
12              $z.p.color = \text{BLACK}$  // case 3
13              $z.p.p.color = \text{RED}$ 
14             RIGHT-ROTATE( $T, z.p.p$ )
15         else (same as then with “right” and “left” exchanged)
16      $T.root.color = \text{BLACK}$ 

```





## Insert fixup loop invariant.

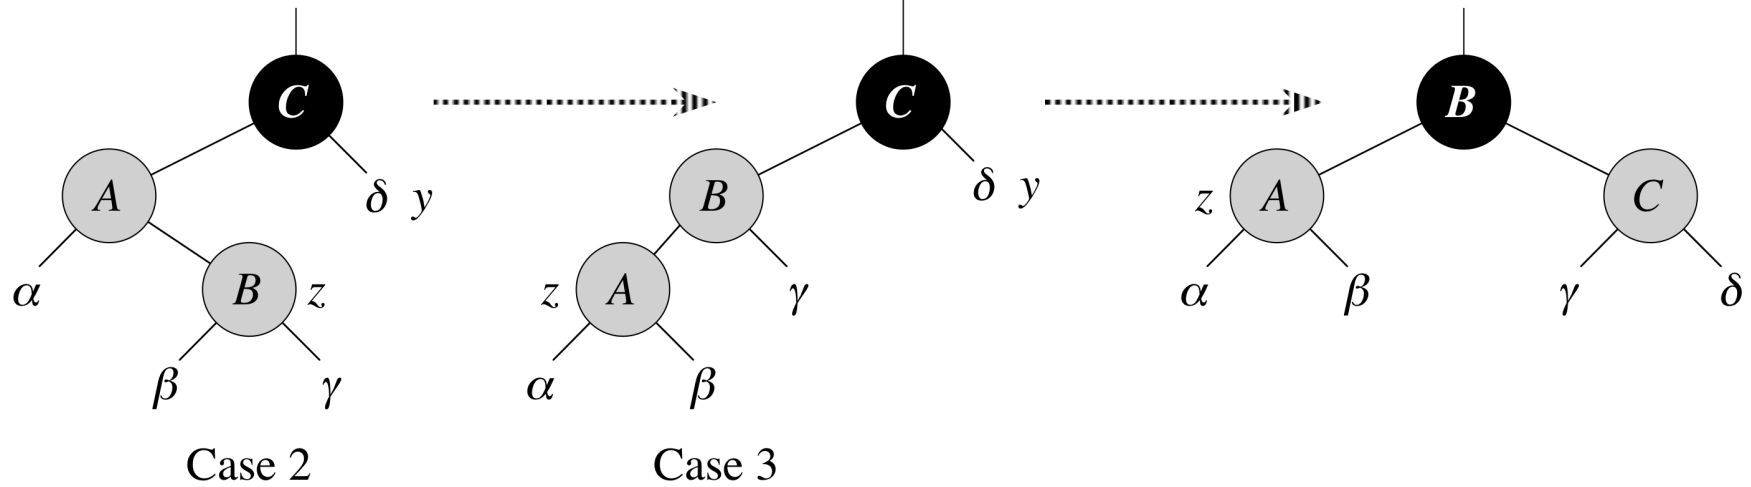
- $z$  is red
- There is at most one red-black violation:
  - $z$  is a red root.
  - $z$  and  $z.p$  are both red.



Parent is red and uncle is red:



Parent is red and uncle is black:



## Analysis

- $O(\lg n)$  time to insert into binary tree.
- Fixup also  $O(\lg n)$ :
  - Each pass through the loop takes  $O(1)$  time.
  - Each iteration moves  $z$  up two levels.
  - $O(\lg n)$  levels.
  - Also note that there are at most 2 rotations overall.
- Insertion into red-black tree is  $O(\lg n)$ .

## Deletion

- Not covered here.
- But also  $O(\lg n)$ .